

Database systems



Wojciech Barczynski

Whoami

Wojciech Barczyński:

- System Engineer
- Lead @ Codility

Quick check

Do you have experience with:

- Programming languages?
- Databases?
- NoSQL?
- Did you ever use SQL?

Aim of the course

1. Relation data model
2. Learn how to query dataases with SQL
3. Design good database
4. Different types of databases

Intuition

- What is a database?
- Let's imagine we need to build one...

Around us

- Youtube
- Self-driving cars

ATM



Operations

- Read Balance
- Update Balance
- Give Money

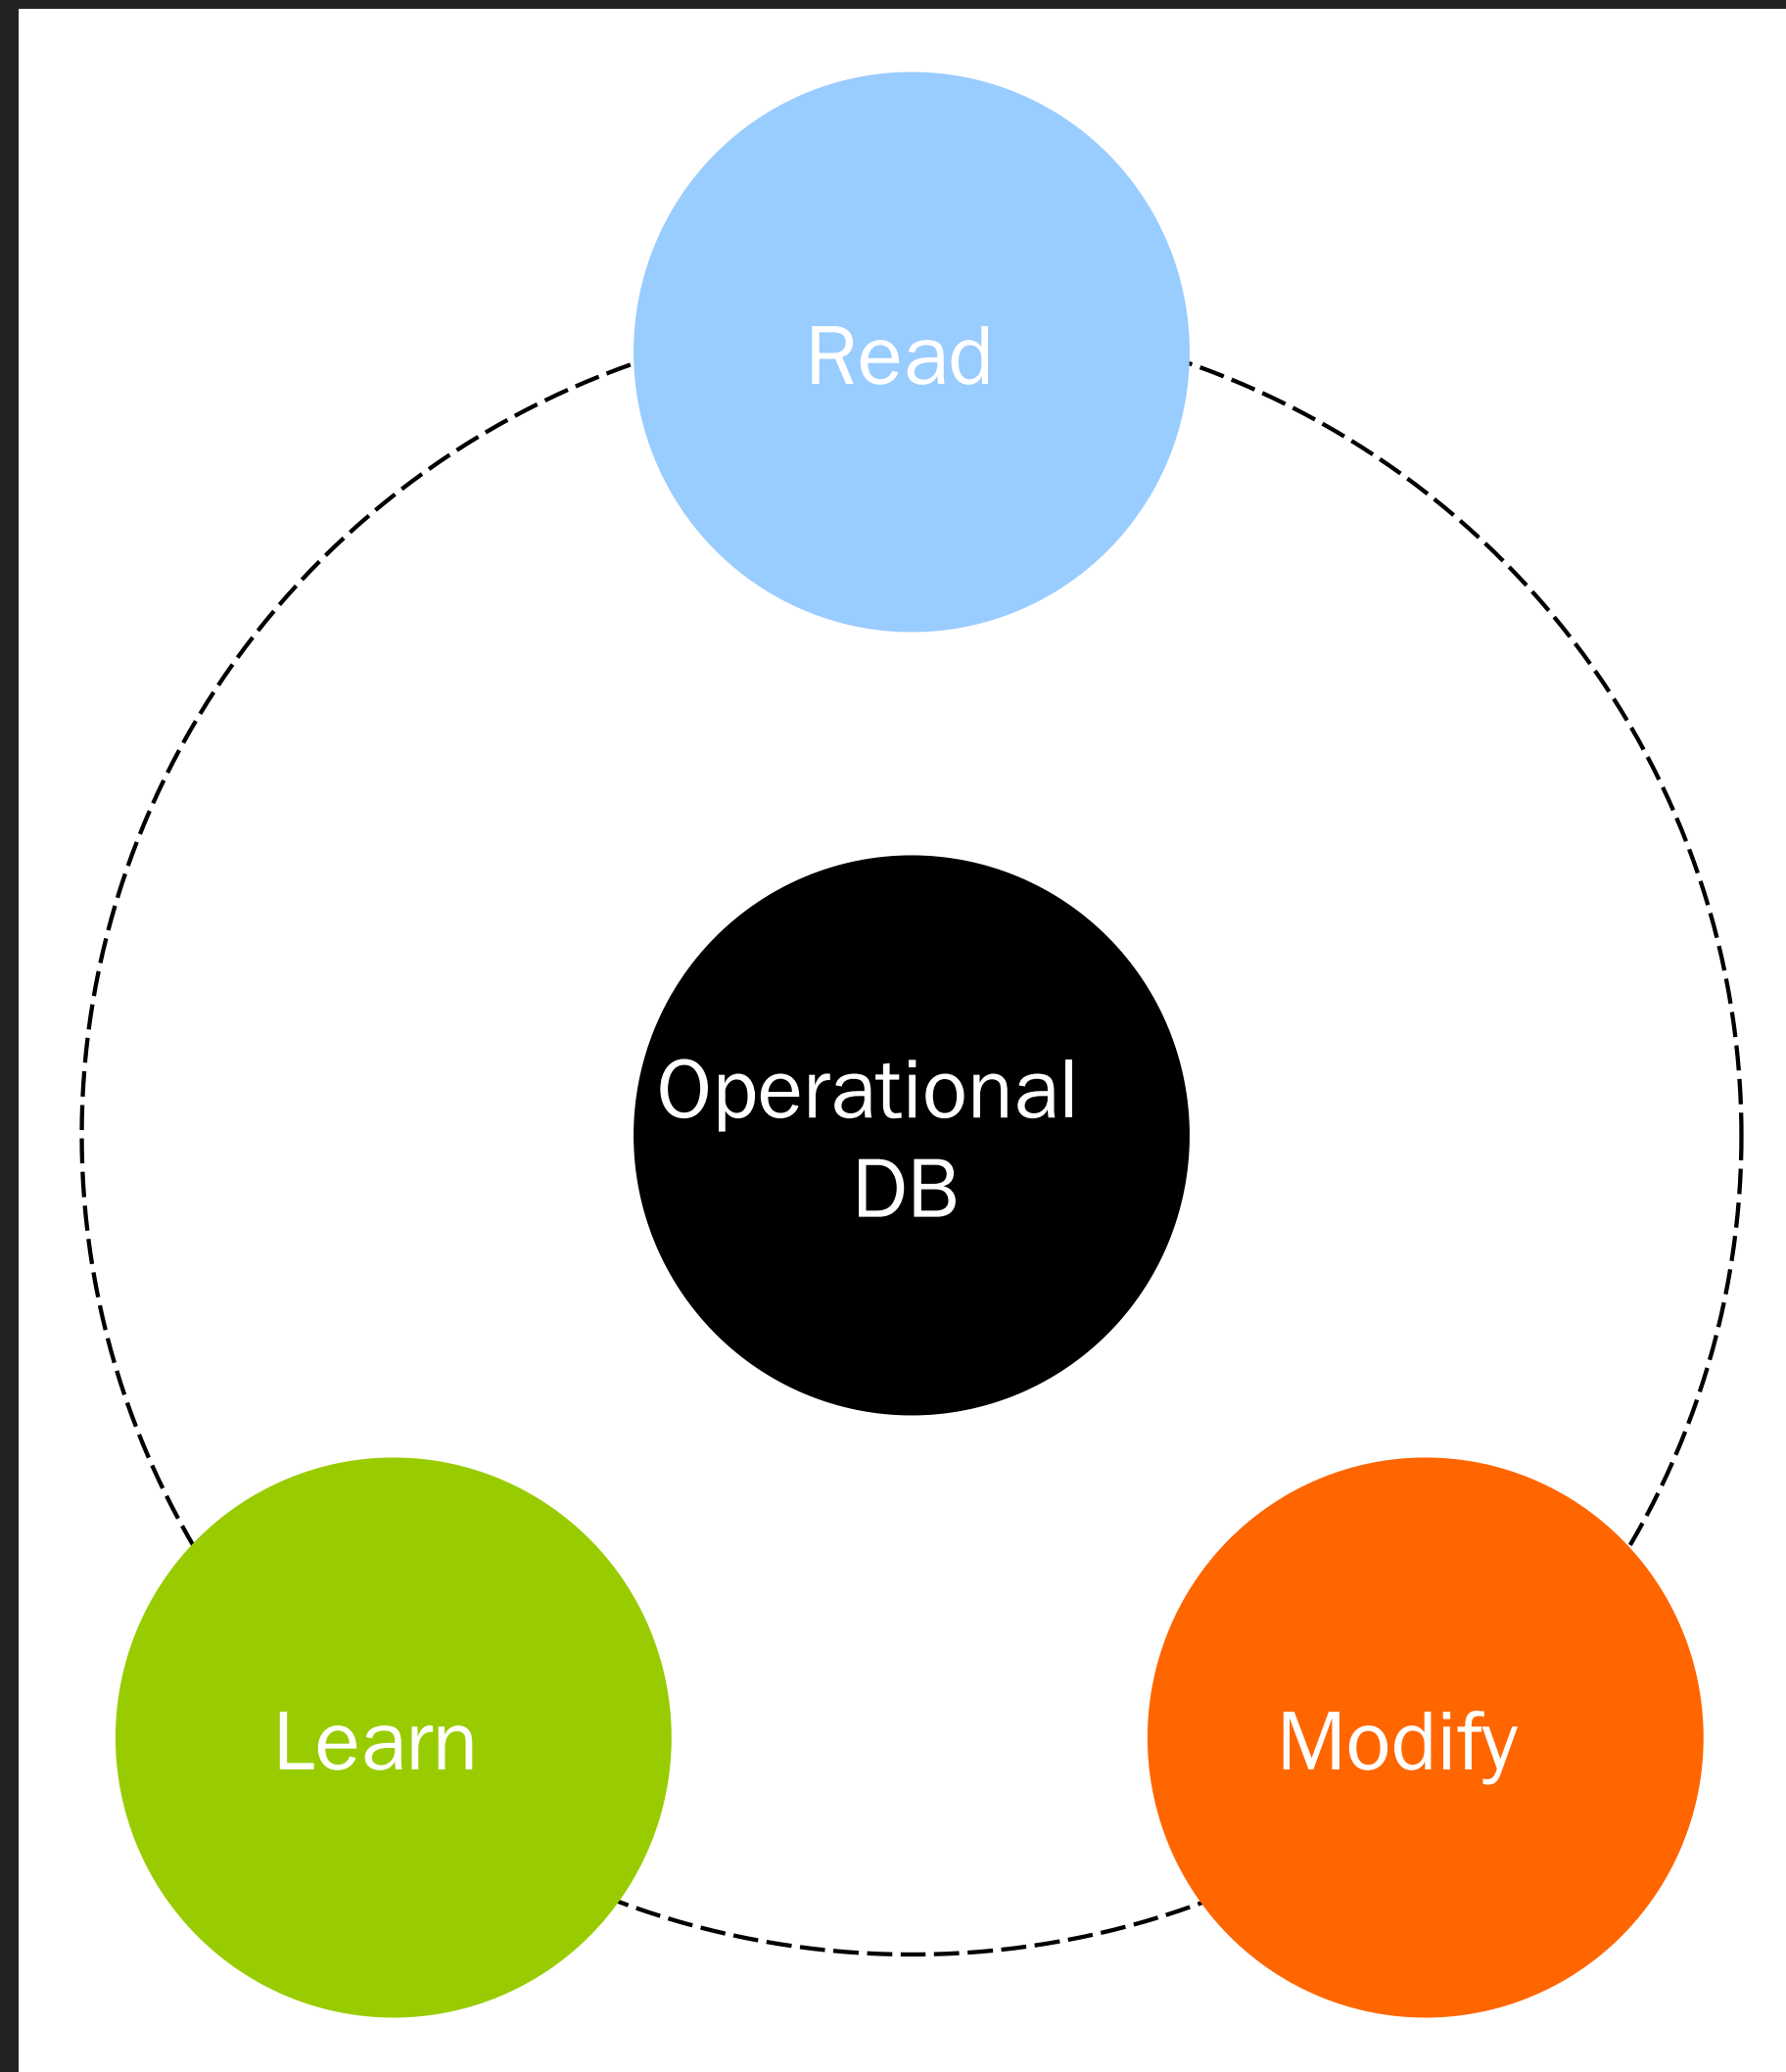
ATM



Operations

- Read Balance
- Give Money
- Update Balance

Standard DB



Supporting

- Scale
- Speed
- Stability
- Reliability

Specialisation

- store data (lots of reads)
- optimize historical data (e.g., logs)
- run batch workflows (e.g., training)

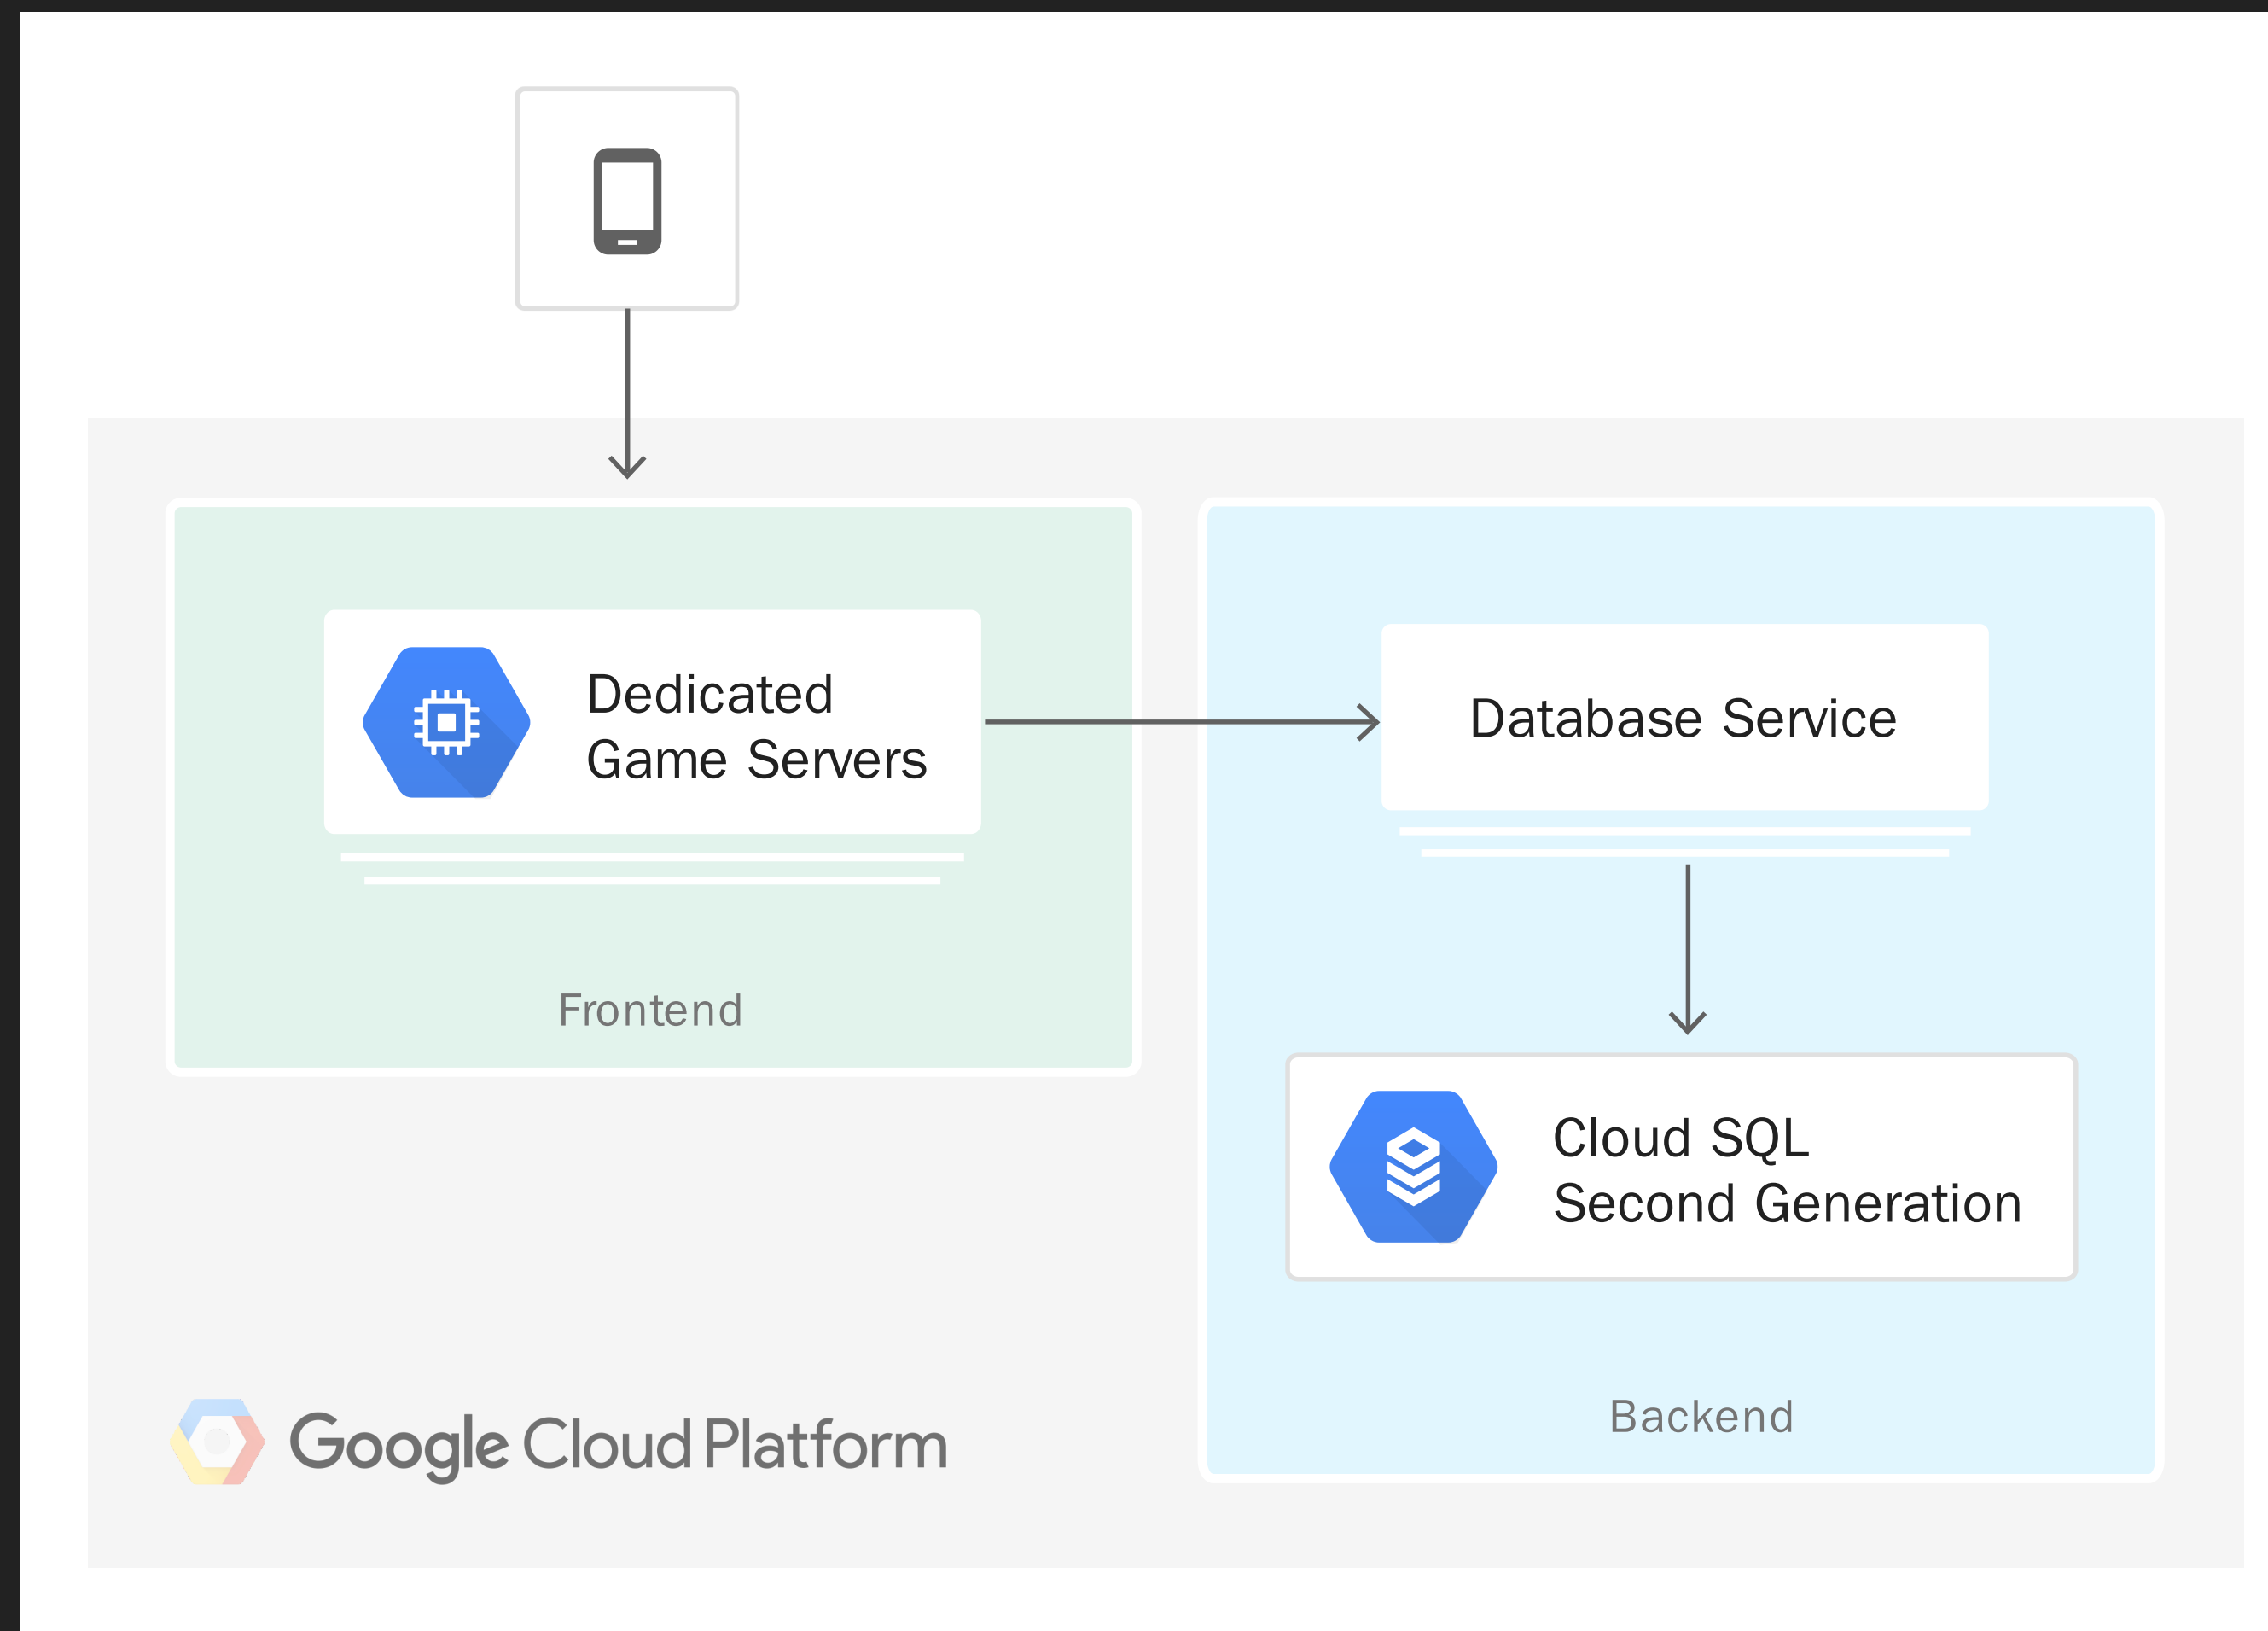
Specialisation

We have over 100+ databases on the market:

- MySQL, Postgres, SAP, Oracle, Sqlite, Mongo,...
- Cloud: GCP, AWS, Bigquery, Coackroach,

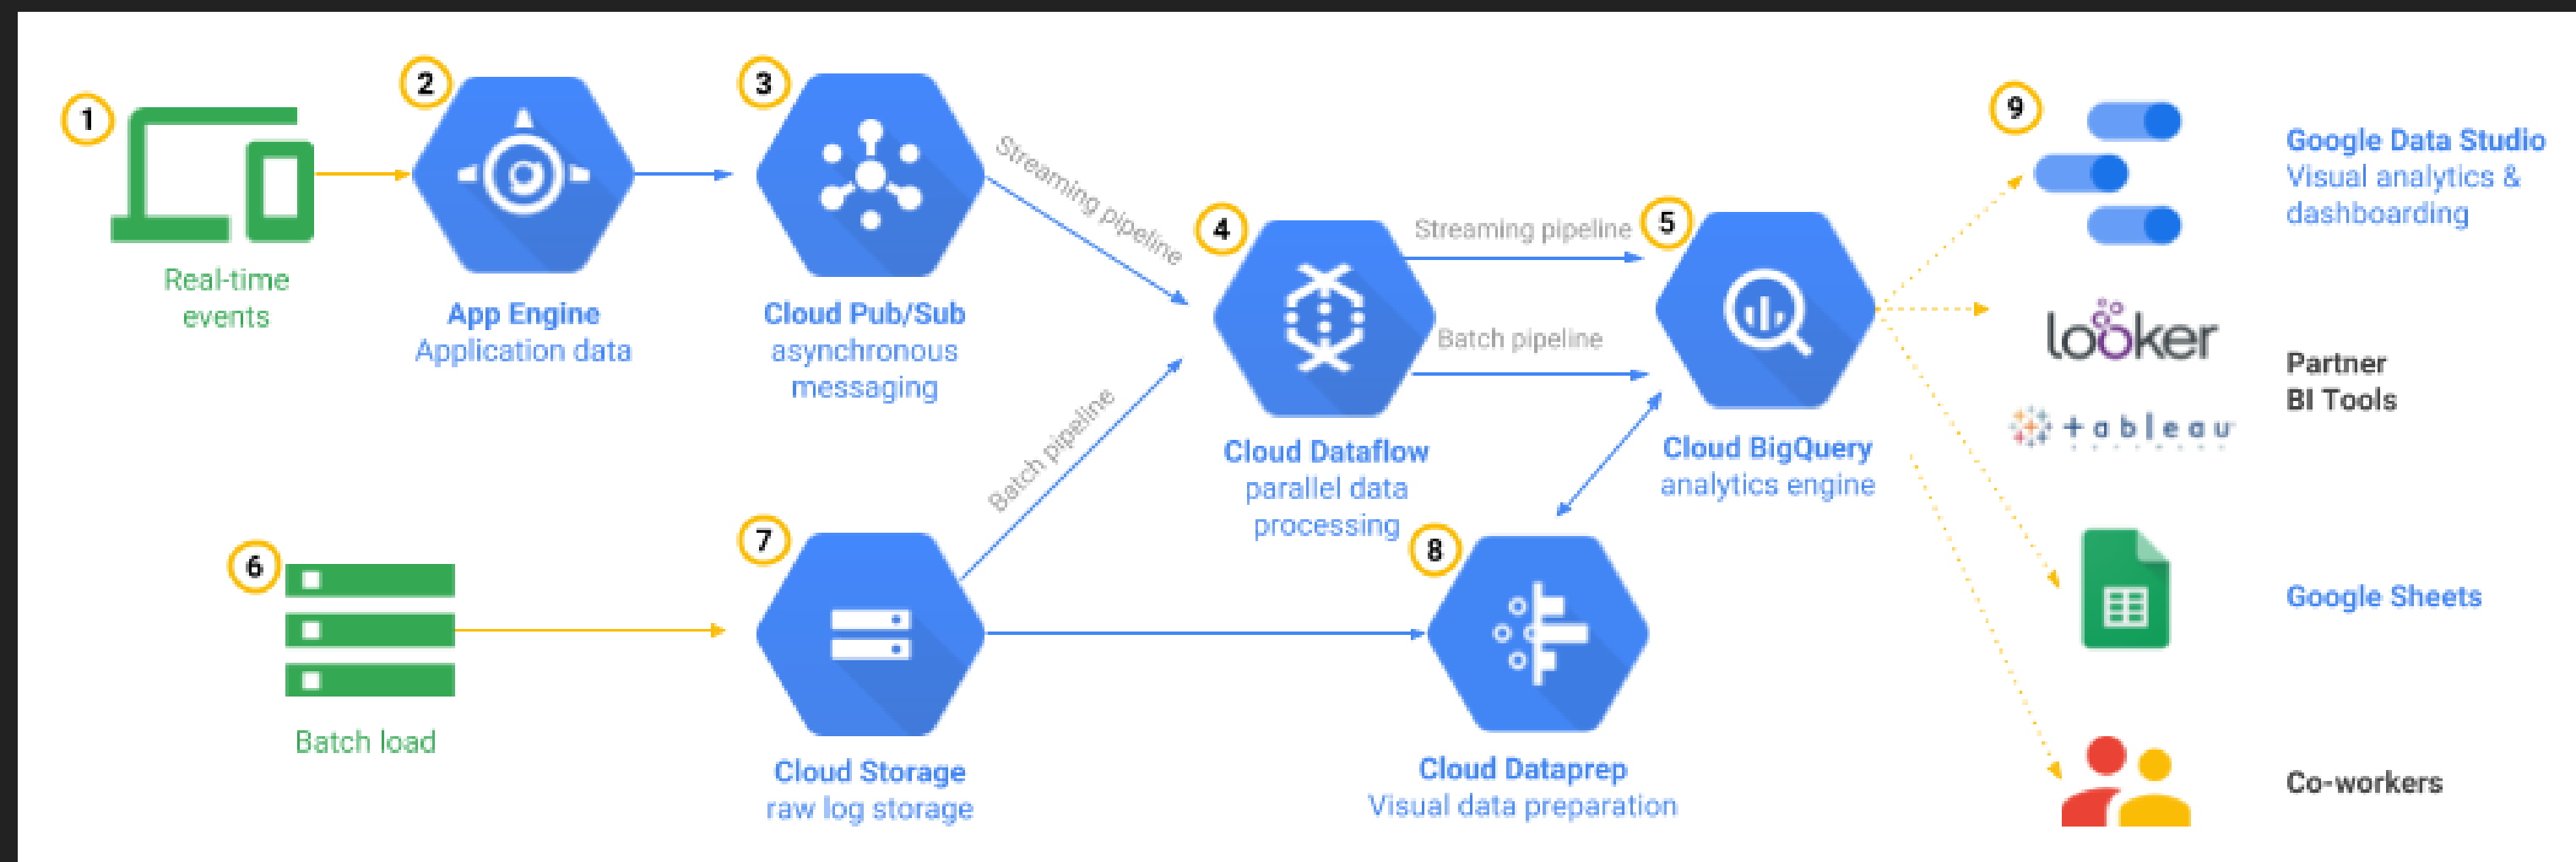
App example

Game app



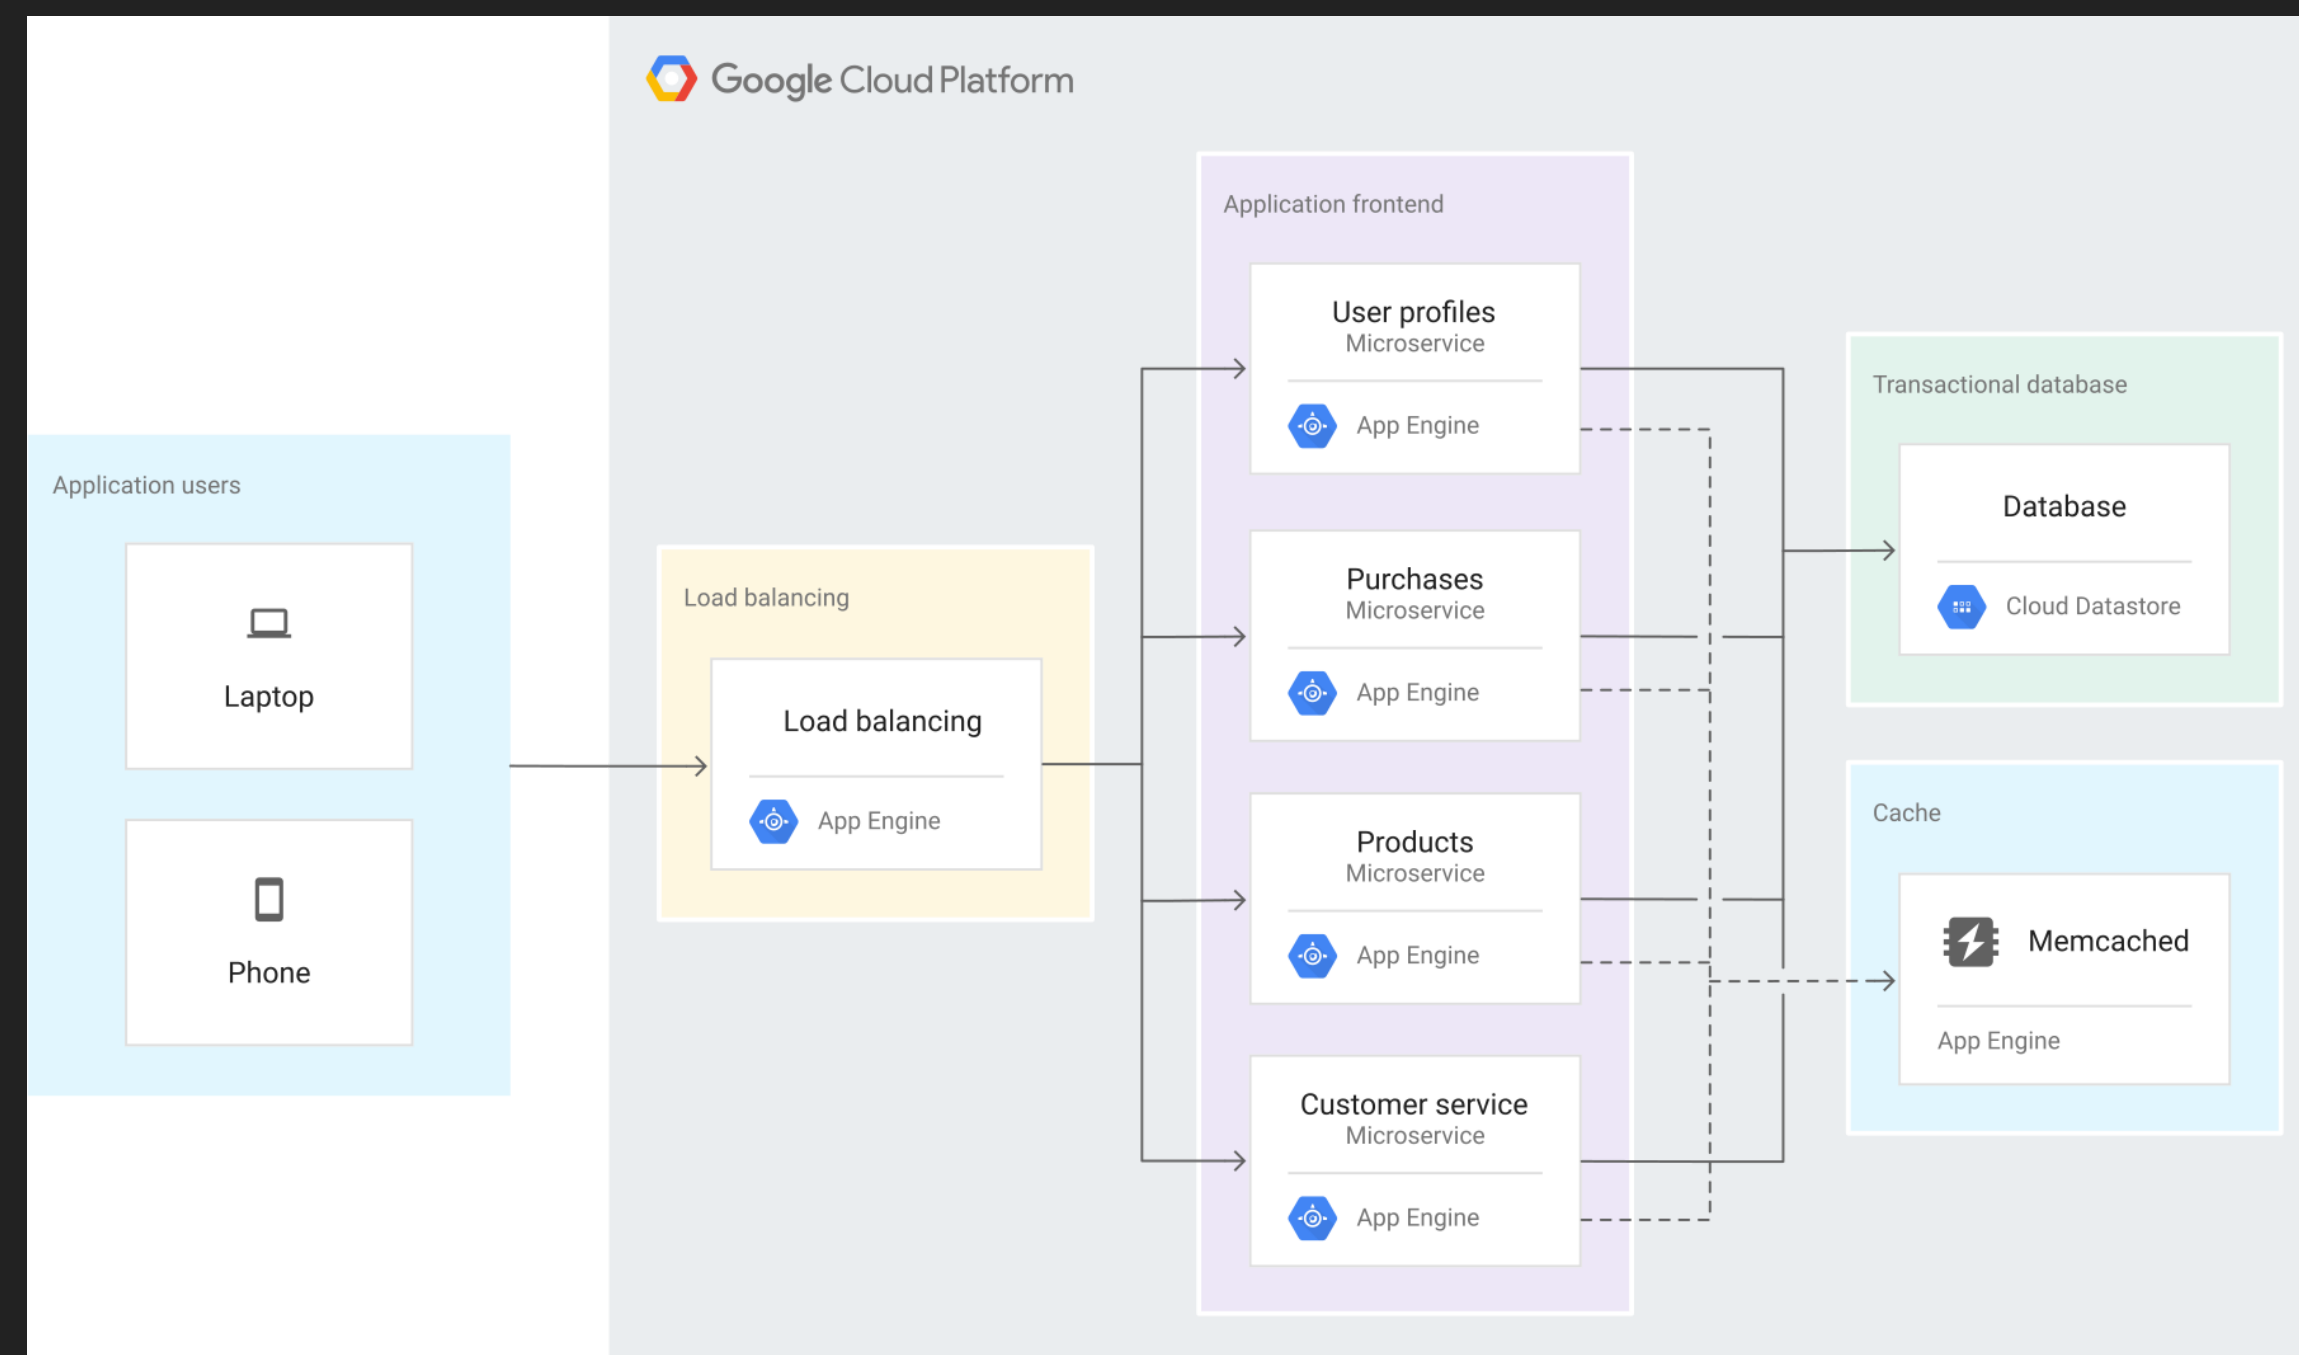
src

Game app



src

Game app



src

Databases

Two Big ideas

- Declarative interfaces
- Transactions

Declarative interfaces

- Apps specify what they want, not how to do it
- Example: “store a table with 2 integer columns”, but not how to encode it on disk
- Example: “count records where $\text{column1} = 5$ ”

Transactions

- Encapsulate multiple app actions into one atomic request (fails or succeeds as a whole)
- Concurrency models for multiple users
- Clear interactions with failure recovery

Declarative interfaces


SQL:

- Abstract “table” data model, many physical implementations
- Specify queries in a restricted language that the database can optimize

Trasaction

- SQL:
 - Commands to start, abort or end transactions
 - based on multiple SQL statements
- Apache Spark: multi-part output of a job appear atomically when all particitns are done

Edgar F Codd




- Proposed the relational DB model, with declarative queries & storage (1970)
- Relation = table with unique key identifying each row

History

Timeline:

- data storage
- navigational databases (1964)
- IBM System R (1974)
- Ingress (1974) - led to PostgreSQL
- Oracle database (1979)

Intro SQL



Intro SQL

	A	B	C	D	E	F	G
1							
2	SID	Name	GPA				
3	3000	Mike	3.7				
4	3001	Joanna	4.8				
5	4002	Tom	5.0				
6	3002	Helene	4.5				
7							
8					SID	CID	Grade
9					3001	1022	5.0
10					3002	1022	3.5
11	CID	Name	Room		4002	1030	4.5
12	1022	Databases	B1				
13	1025	Python	B2				
14	1030	Art	A3				
15							
16							

Tables

- Student(sid:str, name:str, gpa:float)
 - Course(cid:str, name:str, room:str)
 - Enrolled(sid:str, cid:str, grade:float)
- sid → Student
cid → Course

spreadsheet

Intro SQL

	A	B	C	D	E	F	G
1							
2	SID	Name	GPA				
3	3000	Mike	3.7				
4	3001	Joanna	4.8				
5	4002	Tom	5.0				
6	3002	Helene	4.5				
7							
8					SID	CID	Grade
9					3001	1022	5.0
10					3002	1022	3.5
11	CID	Name	Room		4002	1030	4.5
12	1022	Databases	B1				
13	1025	Python	B2				
14	1030	Art	A3				
15							
16							

Queries

- GPA of Student Tom?
- AVG student GPA?
- Helene's classes?

spreadsheet

Example

- [w3schools.com/sql](https://www.w3schools.com/sql)
- `trysql`

SQL

SQL CHEAT SHEET <http://www.sqltutorial.org>



QUERYING DATA FROM A TABLE

SELECT c1, c2 FROM t;
Query data in columns c1, c2 from a table

SELECT * FROM t;
Query all rows and columns from a table

SELECT c1, c2 FROM t
WHERE condition;
Query data and filter rows with a condition

SELECT DISTINCT c1 FROM t
WHERE condition;
Query distinct rows from a table

SELECT c1, c2 FROM t
ORDER BY c1 ASC [DESC];
Sort the result set in ascending or descending order

SELECT c1, c2 FROM t
ORDER BY c1
LIMIT n OFFSET offset;
Skip *offset* of rows and return the next *n* rows

SELECT c1, aggregate(c2)
FROM t
GROUP BY c1;
Group rows using an aggregate function

SELECT c1, aggregate(c2)
FROM t
GROUP BY c1
HAVING condition;
Filter groups using HAVING clause

QUERYING FROM MULTIPLE TABLES

SELECT c1, c2
FROM t1
INNER JOIN t2 ON condition;
Inner join t1 and t2

SELECT c1, c2
FROM t1
LEFT JOIN t2 ON condition;
Left join t1 and t2

SELECT c1, c2
FROM t1
RIGHT JOIN t2 ON condition;
Right join t1 and t2

SELECT c1, c2
FROM t1
FULL OUTER JOIN t2 ON condition;
Perform full outer join

SELECT c1, c2
FROM t1
CROSS JOIN t2;
Produce a Cartesian product of rows in tables

SELECT c1, c2
FROM t1, t2;
Another way to perform cross join

SELECT c1, c2
FROM t1 A
INNER JOIN t2 B ON condition;
Join t1 to itself using INNER JOIN clause

USING SQL OPERATORS

SELECT c1, c2 FROM t1
UNION [ALL]
SELECT c1, c2 FROM t2;
Combine rows from two queries

SELECT c1, c2 FROM t1
INTERSECT
SELECT c1, c2 FROM t2;
Return the intersection of two queries

SELECT c1, c2 FROM t1
MINUS
SELECT c1, c2 FROM t2;
Subtract a result set from another result set

SELECT c1, c2 FROM t1
WHERE c1 [NOT] LIKE pattern;
Query rows using pattern matching %, _

SELECT c1, c2 FROM t
WHERE c1 [NOT] IN value_list;
Query rows in a list

SELECT c1, c2 FROM t
WHERE c1 BETWEEN low AND high;
Query rows between two values

SELECT c1, c2 FROM t
WHERE c1 IS [NOT] NULL;
Check if values in a table is NULL or not

<http://www.sqltutorial.org/sql-cheat-sheet/>

Intro SQL

- slides

Questions?

Wojciech Barczyński

References

- Data Management and Data Systems
- <http://db-class.org/>
- Principles of Data-Intensive Systems

Backup slides

Prepare your env

- Open `Terminal`

```
sudo apt-get install python3-venv python3-pip -y  
sudo snap install atom --classic
```

```
mkdir -p workspace  
cd workspace
```

```
# zauważ kropka  
atom .
```

Prepare your env - Atom

Install the following Atom plugins (Edit->Preferences):

- `platformio-ide-terminal`

Terminal in Atom

- Packages -> platformio-ide-terminal-> Toggle

```
python3 -m venv venv  
source venv/bin/activate
```